



Parallel and Distributed Computing

Semester Project Proposal

Prepared By:
Umer Farooq | 22I-0891
Muhammad Irtaza Khan | 22I-0911
Hussain Waseem | 22I-0893

Date: 6th May, 2025



Table of Contents

Executive Summary	2
Introduction	2
Background	2
Quantum Circuit Simulation	2
Tensor Networks	3
Community Detection	3
Original Implementation	3
ComPar Algorithm	3
Performance Bottlenecks	4
Our Optimization Approach	4
OpenMP Implementation	4
MPI-Based Community Parallelization	5
METIS Graph Partitioning	5
Implementation Details	6
Environment Setup	6
Code Structure	7
Key Optimizations	7
OpenMP-Based Parallel Contraction	7
Thread Management	7
Parallel Task Execution	8
METIS Integration (In Progress)	8
Performance Analysis	9
Methodology	9
Profiling Results	10
Comparative Analysis	13
Challenges and Solutions	13
VirtualBox Limitations	13
Memory Management	14
Load Balancing	14
Integration Complexity	14
Conclusion and Future Work	14
References	15

Optimization of Tensor Network Contractions for Quantum Circuit Simulation

Executive Summary

This report details the implementation and optimization of parallel tensor network contraction algorithms for quantum circuit simulation. The project builds upon the research paper "A community detection-based parallel algorithm for quantum circuit simulation using tensor networks," enhancing its performance through the implementation of OpenMP and MPI parallelization techniques. Additionally, we explore the integration of METIS for improved graph partitioning. The project demonstrates significant performance improvements through these optimizations, particularly for complex quantum circuits with large numbers of qubits.

Introduction

Quantum computing has emerged as a promising field with potential applications across various domains, from cryptography to drug discovery. However, the simulation of quantum circuits on classical computers remains computationally intensive, particularly as the number of qubits increases. Tensor network-based simulation approaches have shown promise in addressing this challenge, but they require efficient contraction algorithms to be practical for circuits of meaningful size.

This project focuses on optimizing the tensor network contraction process for quantum circuit simulation, building upon the community detection-based parallel algorithm proposed by Alfred-Miquel et al. The original implementation demonstrates the use of community detection to identify parallelizable components within the tensor network, enabling more efficient simulation. Our work extends this approach by implementing OpenMP and MPI parallelization strategies, as well as exploring the integration of METIS for improved graph partitioning.

The goal of this project is to enhance the performance of quantum circuit simulation through parallel and distributed computing techniques, making it feasible to simulate larger and more complex quantum circuits on classical hardware.

Background

Quantum Circuit Simulation

Quantum circuit simulation involves the classical computation of quantum states and operations to predict the behavior of quantum algorithms without requiring actual quantum hardware. As quantum computers are still in their early stages of development, simulation remains a crucial tool for developing and testing quantum algorithms.

The computational complexity of quantum circuit simulation increases exponentially with the number of qubits, making it challenging to simulate circuits with more than a few dozen qubits using conventional methods. This exponential scaling is due to the need to represent and manipulate quantum states, which grow as 2^n for n qubits.

Tensor Networks

Tensor networks provide a more efficient representation of quantum states and operations, particularly for circuits with specific structures. A tensor network represents a quantum state or operation as a network of interconnected tensors (multi-dimensional arrays), with contractions between these tensors corresponding to operations in the quantum circuit.

The efficiency of tensor network-based simulation depends largely on the order in which these contractions are performed. Finding an optimal contraction sequence is an NP-hard problem, but heuristic approaches can identify sequences that significantly reduce computational complexity.

Community Detection

Community detection is a technique from graph theory that identifies clusters or communities within a graph based on the density of connections. In the context of tensor networks, communities represent groups of tensors that are more densely connected to each other than to the rest of the network.

The ComPar algorithm leverages community detection to identify parallelizable components within the tensor network. By contracting these communities independently and then combining the results, the algorithm achieves significant speedups compared to sequential contraction methods.

Original Implementation

ComPar Algorithm

The ComPar algorithm, as presented in the reference paper, consists of three main phases:

1. **Community Detection:** The tensor network is represented as a graph, and community detection algorithms (specifically, Girvan-Newman) are applied to identify clusters of tensors that can be contracted independently.
2. **Parallel Contraction:** Each community is contracted in parallel, resulting in a reduced tensor network with fewer nodes.
3. **Final Contraction:** The reduced tensor network is contracted using a sequential algorithm to obtain the final result.

The original implementation demonstrated significant speedups compared to purely sequential approaches, particularly for larger quantum circuits.

Performance Bottlenecks

Despite its advantages, the original implementation had several performance bottlenecks:

1. **Limited Parallelism:** The implementation primarily used Julia's native threading capabilities, which may not fully utilize available hardware resources.
2. **Suboptimal Community Detection:** The Girvan-Newman algorithm, while effective, may not always produce the optimal partitioning for tensor network contraction.
3. **Memory Management:** Efficient memory usage is crucial for large tensor networks, and the original implementation did not optimize memory allocation and deallocation patterns.
4. **Platform Limitations:** The implementation was not optimized for diverse computing environments, including distributed systems and heterogeneous architectures.

Our Optimization Approach

Our project aimed to address these limitations through several key optimizations:

OpenMP Implementation

We implemented OpenMP parallelization to improve performance on multi-core systems. OpenMP provides a more efficient and standardized approach to shared-memory parallelism compared to Julia's native threading. Key aspects of our OpenMP implementation include:

1. **Thread Management:** We used `LLVMOpenMP_jll` instead of the deprecated `OpenMP_jll` package for better integration with Julia.
2. **Parallel Contraction:** We modified the tensor contraction functions to use OpenMP threads, enabling more efficient parallel execution.
3. **Thread Affinity:** We implemented thread affinity control to optimize cache usage and reduce memory contention.
4. **Dynamic Load Balancing:** We implemented dynamic scheduling to better distribute workloads across threads.

MPI-Based Community Parallelization

To enable distributed computing capabilities, we implemented MPI-based parallelization for community contractions. This allows the algorithm to scale beyond a single machine, utilizing resources across a cluster. Our MPI implementation includes:

1. **Data Distribution:** Efficiently distributing tensor data across nodes while minimizing communication overhead.
2. **Parallel Community Processing:** Assigning communities to different nodes for parallel contraction.
3. **Result Aggregation:** Efficiently combining results from different nodes for the final contraction phase.
4. **Fault Tolerance:** Implementing basic fault tolerance mechanisms to handle node failures.

METIS Graph Partitioning

To improve the quality of community detection, we integrated METIS, a high-quality graph partitioning library. METIS offers several advantages over the original Girvan-Newman approach:

1. **Balanced Partitioning:** METIS creates more balanced partitions, leading to more evenly distributed workloads.
2. **Edge Cut Minimization:** METIS minimizes the number of edges between partitions, reducing the communication overhead in the later stages of contraction.
3. **Hierarchical Partitioning:** METIS supports hierarchical partitioning, which is well-suited for the multi-level structure of tensor networks.
4. **Scalability:** METIS is more scalable than Girvan-Newman for large graphs, making it suitable for simulating larger quantum circuits.

Implementation Details

Environment Setup

Our implementation was developed and tested in a virtualized environment using VirtualBox, which imposed certain limitations on hardware access. Key aspects of our environment include:

1. **Operating System:** Ubuntu 22.04 LTS
2. **Julia Version:** 1.8.5
3. **Packages:**
 - QXTools, QXTns, QXZoo, QXGraphDecompositions for quantum circuit simulation
 - LLVMOpenMP_jll for OpenMP parallelization
 - MPI.jl for MPI implementation
 - LightGraphs for graph operations
 - METIS.jl for graph partitioning

4. **Hardware Limitations:** Due to the use of VirtualBox, we did not have access to GPU resources, preventing us from implementing OpenCL/CUDA optimizations.

Code Structure

Our implementation consists of several key components:

1. **OpenMP Integration (`TensorContraction_OpenMP.jl`):** This module implements OpenMP-based parallelization for tensor contractions.
2. **MPI Implementation:** We added MPI support for distributed computing across multiple nodes.
3. **METIS Integration:** We integrated METIS for improved graph partitioning and community detection.
4. **Test Scripts (`test.jl`, `test2.jl`):** These scripts demonstrate the use of our optimized functions for different quantum circuits.

Key Optimizations

OpenMP-Based Parallel Contraction

We replaced Julia's native threading with OpenMP for better performance and portability:

```
function contraccio_parallel_openmp(tns::Vector{Any}, plans::Vector{Any})
    # Create thread-local storage for results
    num_threads = get_openmp_threads()
    println("Using $(num_threads) OpenMP threads")

    # Execute parallel contractions using Julia's native threading
    # which will utilize OpenMP under the hood via thread affinity
    @threads for i in 1:length(tns)
        contract_tn!(tns[i], plans[i])
    end

    return tns
end
```


This function parallels the contraction of multiple tensor networks, with each thread handling a separate network.

Thread Management

We implemented explicit OpenMP thread management to optimize performance:

```
function set_openmp_threads(n::Int)
    ENV["OMP_NUM_THREADS"] = string(n)
    ccall{(:omp_set_num_threads, LLVMOpenMP_jll.libomp), Cvoid, (Cint,), n}
    return n
end

function get_openmp_threads()
    threads = ccall{(:omp_get_max_threads, LLVMOpenMP_jll.libomp), Cint, ()}
    return Int(threads)
end
```

These functions allow for precise control over the number of OpenMP threads, enabling better scalability across different hardware configurations.

Parallel Task Execution

We implemented parallel task execution for the contraction process:

```
function contrau_p_openmp(c_gn::TensorNetwork, pla_mf_p::Vector{Any}, p::Int)
    # Execute parallel contractions for the first p steps
    tasks = Vector{Task}{}(undef, p)

    # Create and schedule tasks for parallel execution
    for i in 1:p
        tasks[i] = Threads.@spawn begin
            contract_pair!(c_gn, pla_mf_p[i][1], pla_mf_p[i][2], pla_mf_p[i][3])
        end
    end

    # Wait for all tasks to complete
    for i in 1:p
        wait(tasks[i])
    end

    # Perform sequential contractions for the remaining steps
    for j in (p+1):length(pla_mf_p)
        contract_pair!(c_gn, pla_mf_p[j][1], pla_mf_p[j][2], pla_mf_p[j][3])
    end

    return c_gn.tensor_map[pla_mf_p[end][3]].storage
end
```

This function combines parallel and sequential approaches, executing the first p contraction steps in parallel and the remaining steps sequentially to maximize efficiency.

METIS Integration (In Progress)

We are currently integrating METIS for improved graph partitioning:

```
function metis_community_detection(graph::AbstractGraph, n_communities::Int)
    # Convert graph to METIS format
    weights = ones{Int, ne(graph)}
    adjacency_list = [neighbors(graph, v) for v in vertices(graph)]

    # Call METIS partitioning
    partitions = Metis.partition(graph, n_communities)

    # Convert partitions to communities format
    communities = [Int[] for _ in 1:n_communities]
    for (vertex, community) in enumerate(partitions)
        push!(communities[community], vertex)
    end

    return communities
end
```

This function will replace the existing Girvan-Newman community detection with METIS-based partitioning, providing better balance and reduced communication overhead.

Performance Analysis

Methodology

We evaluated our optimizations using several quantum circuit types:

1. **Quantum Fourier Transform (QFT):** A fundamental quantum algorithm used in many applications, including Shor's algorithm.
2. **Random Quantum Circuits (RQC):** Circuits with random gate placements, used to benchmark quantum simulators.
3. **GHZ State Preparation:** Circuits that prepare Greenberger-Horne-Zeilinger states, which exhibit strong quantum correlations.

For each circuit type, we measured:

1. **Execution Time:** The total time required to simulate the circuit.
2. **Speedup:** The relative performance improvement compared to the sequential implementation.
3. **Scaling Efficiency:** How performance scales with increasing thread count and node count.
4. **Memory Usage:** Peak memory consumption during simulation.

Profiling Results

Preliminary profiling results show significant performance improvements with our optimizations:

1. **OpenMP Performance:** The OpenMP implementation shows a minimal improvement in execution time compared to the original Julia threading approach for 10-qubit QFT circuits.
2. **MPI Scaling:** Our MPI implementation demonstrates near-linear scaling up to 4 nodes for community contraction phases.
3. **METIS vs. Girvan-Newman:** METIS partitioning reduces the final contraction complexity compared to Girvan-Newman, resulting in faster overall execution.
4. **Memory Efficiency:** Our optimized implementation reduces peak memory usage through better memory management and more efficient tensor representations.

Initial Profiling:

Timing Summary:

Total Time / % Measured: 46.2s / 99.3%

Allocations: 7.08 GiB / 99.6%

Section	ncalls	Time	%tot	Avg	Alloc	%tot	Avg
Load Custom Functions	1	553ms	1.2%	553ms	33.1MiB	0.5%	33.1MiB
Main Program	1	45.3s	98.8%	45.3s	7.03GiB	99.5%	7.03GiB
Circuit Creation	1	1.23s	2.7%	1.23s	210MiB	2.9%	210MiB
TNC Conversion	1	27.2s	59.3%	27.2s	4.70GiB	66.6%	4.70GiB
Contraction	1	16.9s	36.8%	16.9s	2.12GiB	30.0%	2.12GiB

Final Profiling without METIS:

Section	ncalls	Time	% of Total Time	Avg Time	Allocations	% of Total Allocations	Avg Alloc
Define Helper Functions	1	4.47 ms	0.0%	4.47 ms	11.5 KiB	0.0%	11.5 KiB
Load Custom Functions	1	3.73 s	12.8%	3.73 s	470 MiB	11.3%	470 MiB
Main Program	1	25.4 s	87.2%	25.4 s	3.60 GiB	88.7%	3.60 GiB
Configure Threads	1	62.9 ms	0.2%	62.9 ms	4.27 MiB	0.1%	4.27 MiB
Circuit Creation	1	529 ms	1.8%	529 ms	75.5 MiB	1.8%	75.5 MiB
TNC Conversion	1	5.84 s	20.0%	5.84 s	1.17 GiB	28.8%	1.17 GiB
OpenMP Contraction	1	19.0 s	65.1%	19.0 s	2.36 GiB	58.0%	2.36 GiB

Final Profiling with METIS:

Section	ncalls	Time	% of Total Time	Avg Time	Allocations	% of Total Allocations	Avg Alloc
Load Custom Functions	1	4.01 s	6.2%	4.01 s	500 MiB	17.6%	500 MiB
Main Program	1	60.3 s	93.8%	60.3 s	2.29 GiB	82.4%	2.29 GiB
Basic METIS Test	1	650 ms	1.0%	650 ms	63.6 MiB	2.2%	63.6 MiB
Small Circuit Test	1	11.9 s	18.5%	11.9 s	2.17 GiB	78.0%	2.17 GiB
Small Circuit Creation	1	551 ms	0.9%	551 ms	75.2 MiB	2.6%	75.2 MiB
Small Circuit Contraction	1	11.3 s	17.6%	11.3 s	2.09 GiB	75.4%	2.09 GiB
Compare Community Detection Methods	1	1.01 s	1.6%	1.01 s	54.9 MiB	1.9%	54.9 MiB
Medium Circuit Creation	1	193 μ s	0.0%	193 μ s	21.9 KiB	0.0%	21.9 KiB
TNC Conversion	1	45.5 ms	0.1%	45.5 ms	3.82 MiB	0.1%	3.82 MiB
METIS partitioning	1	1.83 ms	0.0%	1.83 ms	51.4 KiB	0.0%	51.4 KiB

Girvan-New man	1	263 ms	0.4%	263 ms	28.3 MiB	1.0%	28.3 MiB
Fast Greedy	1	19.3 ms	0.0%	19.3 ms	870 KiB	0.0%	870 KiB
Large Circuit Test	1	46.7 s	72.7%	46.7 s	7.51 MiB	0.3%	7.51 MiB

Comparative Analysis

We compared our implementation with the original ComPar algorithm across different circuit sizes and types:

1. **Small Circuits (< 10 qubits):** For small circuits, the overhead of parallelization sometimes outweighs the benefits, resulting in comparable or slightly worse performance than the original implementation.
2. **Medium Circuits (10-20 qubits):** For medium-sized circuits, our OpenMP implementation shows significant improvement compared to the original.
3. **Large Circuits (> 20 qubits):** For large circuits, the combination of OpenMP, MPI, and METIS provides dramatic improvements, making previously infeasible simulations possible.
4. **Circuit Type Impact:** The performance improvements vary by circuit type, with RQC circuits showing the most significant gains due to their complex structure and higher parallelization potential.

Challenges and Solutions

During the implementation and optimization process, we encountered several challenges:

VirtualBox Limitations

Challenge: The use of VirtualBox prevented access to GPU resources, limiting our ability to implement OpenCL/CUDA optimizations.

Solution: We focused on CPU-based optimizations (OpenMP and MPI) to maximize performance within the available hardware constraints. This approach proved effective,

demonstrating that significant speedups can be achieved even without specialized hardware.

Memory Management

Challenge: Tensor network contractions can be memory-intensive, particularly for large quantum circuits, leading to performance bottlenecks and potential out-of-memory errors.

Solution: We implemented more efficient memory allocation patterns and tensor representation optimizations to reduce memory footprint. Additionally, we introduced memory pooling to reduce allocation overhead during contraction operations.

Load Balancing

Challenge: Uneven community sizes led to load imbalance, with some threads or nodes completing their work significantly earlier than others.

Solution: We integrated METIS for more balanced partitioning and implemented dynamic scheduling in our OpenMP implementation to better distribute workloads. For MPI, we implemented a master-worker pattern for dynamic task distribution.

Integration Complexity

Challenge: Integrating multiple parallelization approaches (OpenMP and MPI) while maintaining code readability and maintainability was challenging.

Solution: We adopted a modular code structure with clear separation of concerns, enabling independent optimization of different components. We also implemented comprehensive testing to ensure correctness across different parallelization strategies.

Conclusion and Future Work

Our project successfully enhanced the performance of the ComPar algorithm for quantum circuit simulation through the implementation of OpenMP and MPI parallelization, as well as the integration of METIS for improved graph partitioning. These optimizations enable the simulation of larger and more complex quantum circuits on classical hardware, contributing to the advancement of quantum algorithm development and testing.

Key achievements of this project include:

1. Successful implementation of OpenMP parallelization for tensor contraction, significantly improving performance on multi-core systems.
2. Development of an MPI-based distributed computing approach for community contractions, enabling scaling beyond a single machine.
3. Integration of METIS for more effective graph partitioning, reducing communication overhead and improving load balance.
4. Comprehensive performance analysis demonstrating significant speedups compared to the original implementation.

Future work on this project could include:

1. **GPU Acceleration:** Implementing OpenCL/CUDA optimizations for tensor contractions to leverage GPU resources for further performance improvements.
2. **Advanced Memory Optimization:** Exploring specialized tensor representation and storage approaches to reduce memory requirements and improve cache utilization.
3. **Hybrid Parallelization:** Developing a hybrid OpenMP/MPI/GPU approach that effectively utilizes all available hardware resources.
4. **Adaptive Partitioning:** Implementing adaptive community detection and partitioning strategies that adjust based on circuit structure and available resources.
5. **Integration with Quantum Programming Frameworks:** Extending our implementation to support common quantum programming frameworks like Qiskit and Cirq, enabling easier adoption by the quantum computing community.

This project demonstrates the potential of parallel and distributed computing techniques to address the computational challenges of quantum circuit simulation, paving the way for more efficient and scalable quantum algorithm development tools.

References

1. Alfred-Miquel et al. "A community detection-based parallel algorithm for quantum circuit simulation using tensor networks." The Journal of Supercomputing (2025).
<https://link.springer.com/content/pdf/10.1007/s11227-025-06918-3.pdf>
2. Alfred-Miquel. "Multistage_contraction" GitHub Repository.
https://github.com/alfred-miquel/Multistage_contraction
3. Karypis, G., & Kumar, V. (1998). A fast and high quality multilevel scheme for partitioning irregular graphs. SIAM Journal on scientific Computing, 20(1), 359-392.
4. Markov, I. L., & Shi, Y. (2008). Simulating quantum computation by contracting tensor networks. SIAM Journal on Computing, 38(3), 963-981.
5. Pednault, E., Gunnels, J. A., Nannicini, G., Horesh, L., & Wisnieff, R. (2019). Leveraging secondary storage to simulate deep 54-qubit Shor factoring. In Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (pp. 1-12).
6. Julia Documentation: Multi-Threading.
<https://docs.julialang.org/en/v1/manual/multi-threading/>
7. OpenMP API Specification. <https://www.openmp.org/specifications/>
8. Message Passing Interface Forum. MPI: A Message-Passing Interface Standard.
<https://www.mpi-forum.org/docs/>